



UTILITY-BASED AGENTS

Personalized Product Recommendations for E-Commerce

PRESENTED BY:

S.DIVYASRI(192524388)

GUIDED BY:

DR. R. VEERAMANI

What Is a Utility-Based Agent?

A utility-based agent is a rational agent that doesn't just react to its environment — it evaluates every possible action and chooses the one that **maximizes a measurable utility value**.

In e-commerce, "utility" represents how satisfied a customer is likely to be with a recommended product — combining relevance, price fit, quality, and preference alignment into a single comparable score.



Goal-Directed, Not Just Reactive

Unlike simple reflex agents, it weighs trade-offs — e.g., a slightly pricier item with far better reviews — to pick the best overall outcome.

THE AGENT LOOP



Perceive

User behavior, context, catalog data



Evaluate

Score each option by utility



Act

Recommend highest-utility items



Learn

Update model from feedback

Assigning Utility Values to Options

Every candidate product is reduced to one comparable number: a weighted sum of factors that predict customer satisfaction.

$$U(\text{product}) = w_1 \cdot \text{Relevance} + w_2 \cdot \text{Rating} + w_3 \cdot \text{PriceFit} + w_4 \cdot \text{Availability}$$



Relevance

Match to search intent, browsing & purchase history

$w_1 = 0.35$



Rating & Reviews

Aggregate quality signal from other buyers

$w_2 = 0.25$



Price Fit

Alignment with the user's typical spend range

$w_3 = 0.20$




Availability

In stock, fast shipping, delivery reliability

$w_4 = 0.20$

Optimizing at Scale: The Data Pipeline

Millions of users and catalog items mean the agent must score options continuously, using distributed data pipelines rather than one-off calculations.




Real-world scale: systems like these process billions of events per day and re-rank catalogs of tens of millions of items using distributed frameworks (e.g., Spark, Flink) and approximate nearest-neighbor search.

Optimizing for Maximum User Satisfaction

Example: Ranking a Search for "Running Shoes"

Once every candidate has a utility score, the agent simply ranks and recommends the highest scorers — turning subjective satisfaction into an optimizable number.

- ✓ Compares all candidates on one common scale
- ✓ Balances competing goals: price vs. quality vs. speed
- ✓ Re-ranks in real time as new signals arrive
- ✓ Learns from clicks & purchases to refine future weights



Real-World Examples

Utility-based recommendation logic already powers experiences people use every day.



ONLINE RETAIL

"Customers who bought this also bought..."

A large online marketplace scores co-purchase patterns, ratings, and price fit to surface items that maximize the chance of a satisfying purchase.



STREAMING VIDEO

Personalized "For You" rows

A video platform estimates the utility of each title from watch history, completion rate, and genre affinity to rank what appears first.



MUSIC STREAMING

Weekly discovery playlists

A music service blends listening history with collaborative filtering, assigning a utility score to unheard tracks likely to match taste.

CONCLUSION

Why Utility-Based Agents Win



Precision

Reduces recommendations to a single, comparable, optimizable score



Scalability

Distributed pipelines score millions of products for millions of users



Personalization

Weights adapt to individual preference and context, not just popularity



Adaptability

Continuously re-learns from new behavior to stay accurate over time



Takeaway

A utility-based agent turns the vague goal of "make customers happy" into a concrete, data-driven optimization problem — assigning measurable value to every option and using large-scale processing to act on it in real time, at the scale of millions of users and products.